



Why

We model a real-valued output as corrupted by small random errors. Thus, we can talk about a dataset which is “close” to being consistent with a linear predictor.

Definition

Let $(\Omega, \mathcal{A}, \mathbf{P})$ be a probability space. Let $x \in \mathbf{R}^d$ and $e : \Omega \rightarrow \mathbf{R}^n$. For $A \in \mathbf{R}^{n \times d}$, define $y : \Omega \rightarrow \mathbf{R}^n$ by $y = Ax + e$. We call (x, A, e) a *probabilistic errors linear model*. We call y the *response vector*, A the *model matrix* and e the *error vector*.

Moment assumptions

The most basic distributional assumption for a probabilistic errors linear model pertain to the expectation and variance. Since $\mathbf{E}(y) = Ax + \mathbf{E}(e)$ and $\mathbf{var}(y) = \mathbf{var}(e)$, these assumptions can be given for e or for y .

If $\mathbf{E}(x) = 0$ and $\mathbf{var}(y) = \sigma^2 I$ then we call (x, A, e) a *classical linear model with moment assumptions*. Notice that the components of e are assumed uncorrelated. We have $d + 1$ unknowns (the $d \times 1$ entires of θ and scalar parameter σ^2).

In this case $\mathbf{E}(y_i) = a^i{}^\top \theta$ and so θ is called the *mean parameter vector* and σ^2 is called the *model variance*. The model variance indicates the variability inherent in the observations. Neither the mean nor variance of the error depends on the regression vector x nor on the parameter vector θ .

Examples

Consider the *two-sample problem* in which we have two populations with (unknown) mean responses $\alpha_1, \alpha_2 \in \mathbf{R}$. We observe these responses with (perhaps unknown) common variance σ^2 , and assume that errors are uncorrelated.

We define $y^1 = \alpha_1 \mathbf{1} + e^1$ and $y^2 = \alpha_2 \mathbf{1} + e^2$ so that we can stack these and obtain

$$y = \begin{bmatrix} y^1 \\ y^2 \end{bmatrix} = \begin{bmatrix} \alpha_1 \mathbf{1} \\ \alpha_2 \mathbf{1} \end{bmatrix} + \begin{bmatrix} e^1 \\ e^2 \end{bmatrix}.$$

To cast this in our standard form we define

$$A = \begin{bmatrix} \begin{bmatrix} 1 \\ 0 \end{bmatrix} & \cdots & \begin{bmatrix} 1 \\ 0 \end{bmatrix} & \begin{bmatrix} 0 \\ 1 \end{bmatrix} & \cdots & \begin{bmatrix} 0 \\ 1 \end{bmatrix} \end{bmatrix}^\top, \quad x = \begin{bmatrix} \alpha_1 \\ \alpha_2 \end{bmatrix}.$$

with regression vectors $a_1 = (1, 0)$ and $x_2 = (0, 1)$ repeated n_1 and n_2 times, respectively. An input design for this model involves specifying a sequence of these two vectors, which (with the uncorrelated assumption) reduces to dictating how many responses should be collected from each population. The inputs here is really the set $\mathcal{X} = \{1, 2\}$. The feature function is $\phi : \mathcal{X} \rightarrow \mathbf{R}^2$ defined by $\phi(1) = (1, 0)$ and $\phi(2) = (0, 1)$. And so the regression range is $\{(1, 0), (0, 1)\}$.

