



Why

1

Definition

A *variational autoencoder* (VAE) from latent set Z to observation set X is an ordered pair $((p_z^{(\theta)}, p_{x|z}^{(\theta)}), q_{z|x}^{(\phi)})$ whose first coordinate is a deep latent generation pair from Z to X (with parameters θ) and whose second coordinate is deep conditional distribution from X to Z (with parameters ϕ).

A VAE inherits its *joint function* from its deep latent generation pair. $p_z^{(\theta)}$ is called the *latent distribution* (or *prior distribution*, *latent model*). $p_{x|z}^{(\theta)}$ is called the *decoder distribution*. $q_{z|x}^{(\phi)}$ is called the *encoder distribution* (or *inference distribution*, *recognition distribution*).

A *variational autoencoder family*, from Z to X , is a family of autoencoders $\{((p_z^{(\theta)}, p_{x|z}^{(\theta)}), q_{z|x}^{(\phi)})\}_{(\theta, \phi) \in \Theta \times \Phi}$.

¹Future editions will include. Future editions may also change the name of this sheet. It is also likely that there will be added prerequisite sheets on variational inference.

