



Why

We want to summarize a dataset with a distribution.

Overview

The *likelihood* (or *distribution likelihood*) of a probability distribution $p : A \rightarrow \mathbf{R}$ on a dataset $a^1, \dots, a^n \in A$ is $\prod_{i=1}^n p(a^i)$. A *maximum likelihood distribution* $p^* : A \rightarrow \mathbf{R}$ is one which maximizes the likelihood over all distributions on A .

We call the correspondence between datasets and distributions the *maximum likelihood algorithm*. We say that we are selecting the distribution according to the *maximum likelihood principle*. In general, we call any function from datasets to distributions a *distribution selector*.

