



Why

We discuss inferring (or learning) functions from examples.

Definitions

Suppose \mathcal{U} and \mathcal{V} are two sets. A *predictor* from \mathcal{U} to \mathcal{V} is a function $f : \mathcal{U} \rightarrow \mathcal{V}$. We call \mathcal{U} the *inputs*, \mathcal{V} the *outputs*, and $f(u)$ the *prediction* of f on $u \in \mathcal{U}$.

An *inductor* is a function from datasets in $\mathcal{U} \times \mathcal{V}$ to predictors from \mathcal{U} to \mathcal{V} . A *learner* (or *learning algorithm*) is a family of inductors whose index set is \mathbf{N} , and whose n th term is an inductor for a dataset of size n .

Notation

Let D be a dataset of size n in $\mathcal{U} \times \mathcal{V}$. Let $g : \mathcal{U} \rightarrow \mathcal{V}$, a predictor, which makes prediction $g(u)$ on input $u \in \mathcal{U}$. Let $G_n : (\mathcal{U} \times \mathcal{V})^n \rightarrow (\mathcal{U} \times \mathcal{V})$ be an inductor, so that $G_n(D)$ is the predictor which the inductor associates with dataset D . Then $\{G_n : (\mathcal{U} \times \mathcal{V})^n \rightarrow \mathcal{V}^{\mathcal{U}}\}_{n \in \mathbf{N}}$ is a learner.

Relations

Functions are relations, so we might ask if *inferring* relations may be a more general and difficult problem than inferring functions. The following consideration shows that this is *not* the case.

A *relation inductor* is a function from finite datasets in $\mathcal{U} \times \mathcal{V}$ to *relations* on $\mathcal{U} \times \mathcal{V}$. Suppose R is a relation between \mathcal{U} and \mathcal{V} . Suppose the function $f : \mathcal{U} \times \mathcal{V} \rightarrow \{0, 1\}$ is such that

$$f(u, v) = 1 \quad \text{if and only if} \quad (u, v) \in R$$

Given f we can find R , and given R we can find f . Thus, instead of learning the *relation* R we can think of learning the *function* f . In other words, if we have an inductor for f , we have a *relation* inductor for R .

Consistent and complete datasets

What can a dataset tell us?

Suppose $D = ((u_i, v_i))_{i=1}^n$ be a dataset and $R \subset X \times Y$ a relation. D is *consistent with R* if $(u_i, v_i) \in R$ for all $i = 1, \dots, n$. D is *consistent* if there exists a relation with which it is consistent. A dataset is always consistent (take $R = \mathcal{U} \times \mathcal{V}$). D is *functionally consistent* if it is consistent with a function; in this case, $x_i = x_j \Rightarrow y_i = y_j$. D is *functionally complete* if $\cup_i \{x_i\} = X$. In this case, the dataset includes every element of the relation.

Other terminology

Other terms for the inputs include *independent variables*, *explanatory variables*, *precepts*, *covariates*, *patterns*, *instances*, or *observations*. Other terms for the outputs include *dependent variables*, *explained variables*, *postcepts*, *targets*, *outcomes*, *labels* or *observational outcomes*. An input-output pair is sometimes called a *record pair*.

Other terms for a learner include *supervised learning algorithm*. Other terms for a predictor include *input-output mapping*, *prediction rule*, *hypothesis*, *concept*, and *classifier*.

